

Traffic Sign Recognition

Fai Aladhyani

Shahad Alotaibi

Maram Alshehri

Abstract

This project applies deep learning to a traffic-related problem, specifically Traffic Sign Recognition. We designed and implemented a model using a chosen dataset. Our results demonstrate the potential of deep learning in solving traffic-related problems and highlight the importance of careful data preprocessing and model selection

Table of Contents

- Abstract
- Introduction
- Objectives
- Conclusions
- Our Team
- List of figures
- Questions

Introduction

As part of the Project: Traffic-Related Deep Learning Project, our objective is to design and implement a deep learning model that can effectively solve a traffic-related problem. We have chosen to focus on Traffic sign recognition, a critical aspect of traffic management that has significant implications for road safety and efficiency.

Objective

The objectives of this project are:

- To select a traffic-related problem and implement a deep learning model to address it.
- To collect and preprocess a dataset relevant to the chosen problem.
- To design and implement a suitable deep learning model using a chosen framework (e.g., TensorFlow, PyTorch).
- To train and evaluate the model using relevant performance metrics.
- To fine-tune hyperparameters to optimize model performance.
- To document the entire workflow, including data collection, preprocessing, model selection, implementation, training, and evaluation.
- To present the project's findings, including results, challenges, and conclusions

Our Team :

Name	Tasks
Shahad Alotaibi	Collecting data, Code, trying Experiments on our model, edit presentation.
Maram Alshehri	Collecting data, Code, trying Experiments on our model, prepare presentation.
Fai Aladhyani	Collecting data, trying Experiments on our model, Write the Report

Conclusion

Through this project, we gained hands-on experience with deep learning frameworks and tools, including TensorFlow, PyTorch, Keras ,We also developed our skills in data preprocessing, model design, and evaluation

List All Figures

Traffic Sign Recognition using Deep Learning(CNN)

Traffic sign recognition using deep learning leverages Convolutional Neural Networks (CNNs) to classify traffic signs from images. CNNs automatically extract features from images and learn spatial hierarchies, making them highly effective for image classification tasks.

Import libraries

This script outlines a basic image classification workflow using Keras for deep learning. It includes data processing, model building, and model training.

```
[9]:  
import os  
import cv2  
import numpy as np  
import pickle  
import random  
import numpy as np  
import seaborn as sns  
import pandas as pd  
import matplotlib.pyplot as plt  
from keras.models import Sequential  
from keras.layers import Dense  
from keras.optimizers import Adam  
from keras.losses import categorical_crossentropy  
from keras.utils import to_categorical  
from keras.layers import Dropout, Flatten  
from keras.layers import Conv2D, MaxPooling2D  
from sklearn.model_selection import train_test_split  
from tensorflow.keras.preprocessing.image import ImageDataGenerator  
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau, ModelCheckpoint  
from sklearn.metrics import classification_report
```

Configuration Variables for Traffic Sign Recognition

This section defines various configuration variables used for setting up and managing the traffic sign recognition project using deep learning.

```
[7]:  
path = "Dataset"  
labelFile = 'labels.csv'  
batch_size_val=32  
epochs_val=20  
imageDimensions = (32,32,3)  
testRatio = 0.2  
validationRatio = 0.2
```

```

# Get the predicted labels
y_pred = model.predict(X_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_true = np.argmax(y_test, axis=1)

# Generate and print classification report
report = classification_report(y_true, y_pred_classes)
print("Classification Report:")
print(report)

```

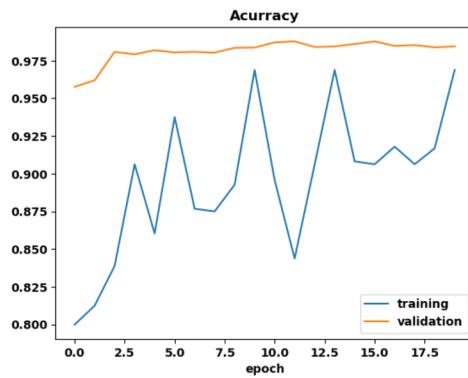
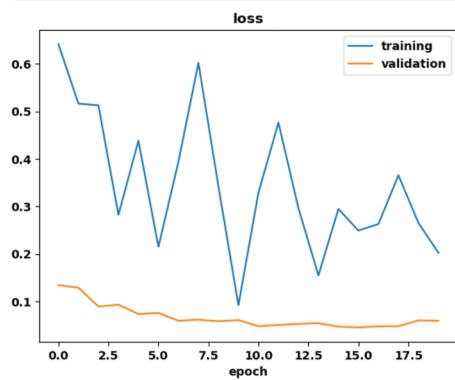
218/218 ————— 7s 3ims/step

	precision	recall	f1-score	support
0	1.00	0.95	0.97	39
1	0.98	1.00	0.99	354
2	0.99	0.98	0.99	394
3	0.99	0.96	0.98	230
4	0.98	0.99	0.98	326
5	0.96	0.97	0.97	339
6	0.99	1.00	0.99	70
7	0.91	0.99	0.95	257
8	1.00	0.91	0.95	263
9	1.00	0.99	0.99	264
10	1.00	1.00	1.00	370
11	0.99	1.00	0.99	209
12	0.99	1.00	0.99	364
13	1.00	1.00	1.00	384
14	1.00	0.99	1.00	132
15	0.99	0.96	0.98	110
16	1.00	1.00	1.00	65
17	1.00	1.00	1.00	208
18	0.94	0.99	0.97	241
19	1.00	0.81	0.89	31
20	1.00	0.99	0.99	68
21	0.94	1.00	0.97	48
22	1.00	1.00	1.00	74
23	1.00	0.97	0.99	69
24	0.98	1.00	0.99	46
25	1.00	1.00	1.00	293
26	0.94	0.87	0.90	100
27	0.94	0.98	0.96	46
28	0.92	0.99	0.96	98
29	0.96	0.85	0.90	53
30	0.99	0.95	0.97	85
31	0.99	0.98	0.99	139
32	1.00	1.00	1.00	42
33	1.00	0.98	0.99	131
34	1.00	1.00	1.00	92
35	1.00	0.99	1.00	248
36	0.97	1.00	0.98	62
37	1.00	1.00	1.00	32
38	0.99	1.00	1.00	387
39	0.98	0.98	0.98	54
40	1.00	0.99	0.99	67
41	1.00	0.98	0.99	42
42	1.00	0.97	0.99	34
accuracy			0.98	6960
macro avg	0.98	0.98	0.98	6960
weighted avg	0.98	0.98	0.98	6960

```

1 [...]
    plt.figure(1)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.legend(['training', 'validation'])
    plt.title('loss')
    plt.xlabel('epoch')
    plt.figure(2)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.legend(['training', 'validation'])
    plt.title('Accuracy')
    plt.xlabel('epoch')
    plt.show()

```



Model Evaluation:

```

1 [...]
    score = model.evaluate(X_test,y_test,verbose=0)
    print('Test Loss:',score[0])
    print('Test Accuracy:',score[1])

```

Test Loss: 0.06176169216632843
 Test Accuracy: 0.9833333492279053

```

...
history=model.fit(dataGen.flow(X_train,y_train,batch_size=batch_size_val),steps_per
Epoch 1/20
695/695 68s 98ms/step - accuracy: 0.7891 - loss: 0.6722 - val
_accuracy: 0.9576 - val_loss: 0.1340
Epoch 2/20
695/695 5s 7ms/step - accuracy: 0.8125 - loss: 0.5161 - val_a
ccuracy: 0.9619 - val_loss: 0.1287
Epoch 3/20
695/695 65s 94ms/step - accuracy: 0.8369 - loss: 0.5256 - val
_accuracy: 0.9808 - val_loss: 0.0890
Epoch 4/20
695/695 5s 6ms/step - accuracy: 0.9062 - loss: 0.2819 - val_a
ccuracy: 0.9792 - val_loss: 0.0929
Epoch 5/20
695/695 67s 96ms/step - accuracy: 0.8526 - loss: 0.4613 - val
_accuracy: 0.9819 - val_loss: 0.0733
Epoch 6/20
695/695 4s 6ms/step - accuracy: 0.9375 - loss: 0.2146 - val_a
ccuracy: 0.9804 - val_loss: 0.0754
Epoch 7/20
695/695 66s 94ms/step - accuracy: 0.8779 - loss: 0.3991 - val
_accuracy: 0.9808 - val_loss: 0.0592
Epoch 8/20
695/695 4s 6ms/step - accuracy: 0.8750 - loss: 0.6019 - val_a
ccuracy: 0.9802 - val_loss: 0.0613
Epoch 9/20
695/695 66s 94ms/step - accuracy: 0.8931 - loss: 0.3404 - val
_accuracy: 0.9835 - val_loss: 0.0581
Epoch 10/20
695/695 4s 6ms/step - accuracy: 0.9688 - loss: 0.0918 - val_a
ccuracy: 0.9837 - val_loss: 0.0605
Epoch 11/20
695/695 77s 111ms/step - accuracy: 0.8934 - loss: 0.3355 - va
l_accuracy: 0.9871 - val_loss: 0.0476
Epoch 12/20
695/695 7s 9ms/step - accuracy: 0.8438 - loss: 0.4760 - val_a
ccuracy: 0.9878 - val_loss: 0.0502
Epoch 13/20
695/695 63s 90ms/step - accuracy: 0.9076 - loss: 0.2897 - val
_accuracy: 0.9840 - val_loss: 0.0524
Epoch 14/20
695/695 4s 6ms/step - accuracy: 0.9688 - loss: 0.1546 - val_a
ccuracy: 0.9844 - val_loss: 0.0540
Epoch 15/20
695/695 66s 95ms/step - accuracy: 0.9084 - loss: 0.2919 - val
_accuracy: 0.9860 - val_loss: 0.0467
Epoch 16/20
695/695 4s 6ms/step - accuracy: 0.9062 - loss: 0.2489 - val_a
ccuracy: 0.9878 - val_loss: 0.0451
Epoch 17/20
695/695 67s 96ms/step - accuracy: 0.9183 - loss: 0.2575 - val
_accuracy: 0.9847 - val_loss: 0.0474
Epoch 18/20
695/695 5s 7ms/step - accuracy: 0.9062 - loss: 0.3653 - val_a
ccuracy: 0.9853 - val_loss: 0.0475
Epoch 19/20
695/695 76s 109ms/step - accuracy: 0.9177 - loss: 0.2689 - va
l_accuracy: 0.9838 - val_loss: 0.0599
Epoch 20/20
695/695 4s 6ms/step - accuracy: 0.9688 - loss: 0.2022 - val_a
ccuracy: 0.9844 - val_loss: 0.0592

```

```

print(model.summary())
model: "sequential"



| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d_1 (Conv2D)<br>display   | (None, 28, 28, 60) | 1,560   |
| conv2d_1 (Conv2D)              | (None, 24, 24, 60) | 90,060  |
| max_pooling2d (MaxPooling2D)   | (None, 12, 12, 60) | 0       |
| conv2d_2 (Conv2D)              | (None, 10, 10, 30) | 16,230  |
| conv2d_3 (Conv2D)              | (None, 8, 8, 30)   | 8,130   |
| max_pooling2d_1 (MaxPooling2D) | (None, 4, 4, 30)   | 0       |
| dropout (Dropout)              | (None, 4, 4, 30)   | 0       |
| flatten (Flatten)              | (None, 480)        | 0       |
| dense (Dense)                  | (None, 500)        | 240,500 |
| dropout_1 (Dropout)            | (None, 500)        | 0       |
| dense_1 (Dense)                | (None, 43)         | 21,543  |

Total params: 378,023 (1.44 MB)
Trainable params: 378,023 (1.44 MB)
Non-trainable params: 0 (0.00 B)
None

```

Data Augmentation:

```
dataGen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, zoom_range=0.2, shear_range=0.1, rotation_range=0.1, fill_mode='nearest', c
```

```
dataGen.fit(X_train)
batches = dataGen.flow(X_train,y_train,batch_size=20)
X_batch,y_batch = next(batches)
```

Converting Labels to Categorical Format:

```
y_train = to_categorical(y_train,noOfClasses)
y_validation = to_categorical(y_validation,noOfClasses)
y_test = to_categorical(y_test,noOfClasses)
```

CNN Model Creation:

This code creates, compiles, trains, and evaluates a Convolutional Neural Network (CNN) model using TensorFlow/Keras. The model is designed for image classification, likely for tasks such as Traffic Sign Recognition, as suggested by the previous context. The code also includes visualizations of training progress and evaluation on the test set.

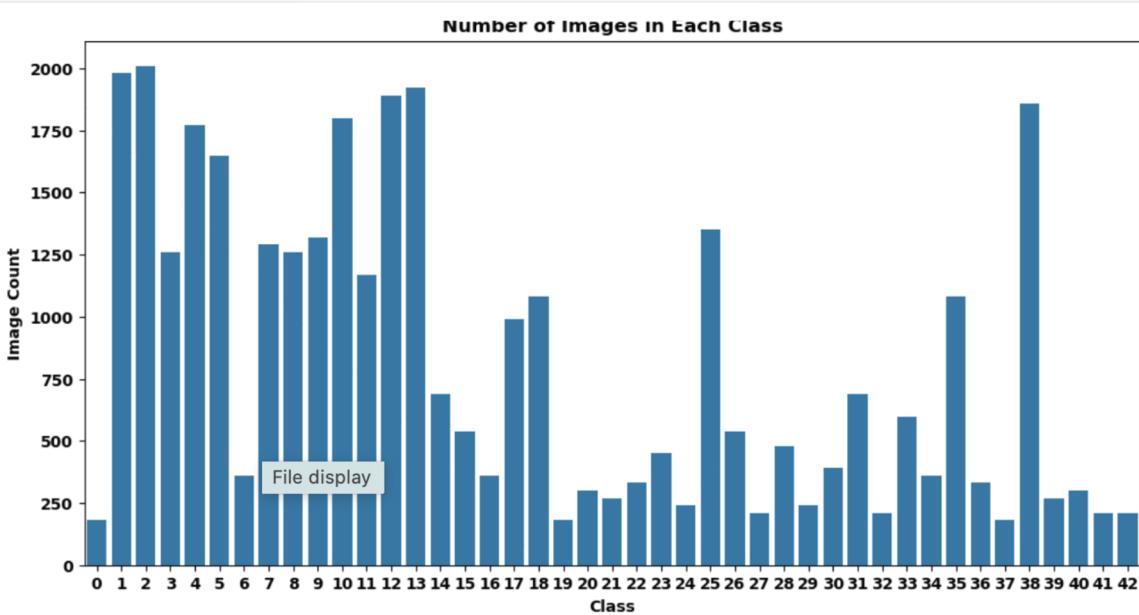
```
model= Sequential()
model.add((Conv2D(60,(5,5),input_shape=(imageDimensions[0],imageDimensions[1],1),activation='relu'))) # ADDING MORE LAYERS
model.add((Conv2D(60, (5,5), activation='relu')))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add((Conv2D(30, (3,3),activation='relu')))
model.add((Conv2D(30, (3,3), activation='relu')))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.5))

model.add(Flatten())
model.add(Dense(500,activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(noOfClasses,activation='softmax'))
```

c:\Users\HP\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass a n `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
model.compile(optimizer=Adam(learning_rate=0.001), loss=categorical_crossentropy, metrics=[ 'accuracy' ])
```



Initializing Preprocessing Functions:

```

def grayscale(img):
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    return img
def equalize(img):
    img = cv2.equalizeHist(img)
    return img
def preprocessing(img):
    img = grayscale(img)
    img = equalize(img)
    img = img/255
    return img

```

Preprocessing Training, Validation, and Test Sets:

```

X_train=np.array(list(map(preprocessing,X_train)))
X_validation=np.array(list(map(preprocessing,X_validation)))
X_test=np.array(list(map(preprocessing,X_test)))

```

Reshaping the Data:

```

X_train=X_train.reshape(X_train.shape[0],X_train.shape[1],X_train.shape[2],1)
X_validation=X_validation.reshape(X_validation.shape[0],X_validation.shape[1],X_validation.shape[2],1)
X_test=X_test.reshape(X_test.shape[0],X_test.shape[1],X_test.shape[2],1)

```

Printing the Shapes of Train, Validation, and Test Data

This block of code is used to display the shapes (dimensions) of the training, validation, and testing datasets for both features (images) and labels (class numbers). It helps ensure that the data has been split correctly into the respective sets.

```
print("Data Shapes")
print("Train",end = "");print(X_train.shape,y_train.shape)
print("Validation",end = "");print(X_validation.shape,y_validation.shape)
print("Test",end = "");print(X_test.shape,y_test.shape)
```

```
Data Shapes
Train(22271, 32, 32, 3) (22271,)
Validation(5568, 32, 32, 3) (5568,)
Test(6960, 32, 32, 3) (6960,)
```

Data Preprocessing and Augmentation for Training, Validation, and Test Data

Loading Label Data:

```
data=pd.read_csv(labelFile)
print("data shape ",data.shape,type(data))

data shape  (43, 2) <class 'pandas.core.frame.DataFrame'>
```

```
import os

# Path to the directory containing the class folders
dataset_dir = "Dataset"

# Initialize a dictionary to store the count of images for each class
class_image_count = {}

# Iterate over the class folders (assuming they are named from 0 to 42)
for class_name in range(43): # Assuming class names are from 0 to 42
    class_dir = os.path.join(dataset_dir, str(class_name))
    if os.path.exists(class_dir) and os.path.isdir(class_dir):
        # Count the number of files (images) in the class directory
        num_images = len(os.listdir(class_dir))
        class_image_count[class_name] = num_images

# Create a DataFrame from the class_image_count dictionary
df = pd.DataFrame(list(class_image_count.items()), columns=['Class', 'Image Count'])

# Plotting the countplot using Seaborn
plt.figure(figsize=(12, 6))
sns.barplot(data=df, x='Class', y='Image Count')
plt.title('Number of Images in Each Class')
plt.xlabel('Class')
plt.ylabel('Image Count')
plt.show()
```



Data Loading and Preparation

This section of the code handles loading and preparing the dataset of traffic signs for further processing and model training.

```
count = 0
images = []
classNo = []
myList = os.listdir(path)
print("Total Classes Detected:", len(myList))
noOfClasses=len(myList)
print("Importing Classes.....")
for x in range (0,len(myList)):
    myPicList = os.listdir(path+"/"+str(count))
    for y in myPicList:
        curImg = cv2.imread(path+"/"+str(count)+"/"+y)
        images.append(curImg)
        classNo.append(count)
        print(count, end = " ")
        count +=1
print(" ")
images = np.array(images)
classNo = np.array(classNo)

Total Classes Detected: 43
Importing Classes.....
```

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41
42

Data Splitting for Training, Testing, and Validation

```
X_train, X_test, y_train, y_test = train_test_split(images, classNo, test_size=testRatio)
X_train, X_validation, y_train, y_validation = train_test_split(X_train, y_train, test_size=validationRatio)
```

Displaying Random Images from the Dataset

In this block of code, we are displaying a grid of 25 random images from the combined training and testing datasets.

```
# Combine the X_train and X_test sets for displaying images
all_images = np.concatenate((X_train, X_test))

# Choose 25 random images
random_images = random.sample(range(all_images.shape[0]), 25)

# Display the random images
plt.figure(figsize=(10, 10))
for i, idx in enumerate(random_images):
    image = all_images[idx]

    plt.subplot(5, 5, i+1)
    plt.imshow(image)
    plt.axis('off')
```

Model Saving:

```
[...  
    model.save("model.h5")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Conclusion

The CNN model achieved strong performance on the image classification task through effective preproc improvements could focus on fine-tuning the model or exploring deeper architectures for better accuracy.

Questions

What challenges did you face during data collection and how did you overcome them?

Data Collecting took a lot of the time to find the appropriate data size and of our target.

Why did you choose your specific model architecture?

CNN because it is good with object recognition.

How does the performance of your chosen model compare to alternative approaches, and what factors contributed to the differences?

We did not try another model to compare, we only try and used CNN because for time constrain.

What did you learn from storing the data in a database, and how did it impact your Workflow?

We did not used a Database.

If you were to extend this project, what additional features or improvements would you consider?

Future improvements:

- Experiment with over-sampling and under-sampling techniques
- Add more signs.
- Improve model accuracy.
- Real-time traffic sign detection